

# Customer Churn

April 3, 2025

## 1 Customer Churn Prediction

```
[73]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.ensemble import RandomForestClassifier
import statsmodels.api as sm
from sklearn import metrics
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.metrics import accuracy_score, classification_report, log_loss
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc
from sklearn.pipeline import Pipeline
from sklearn.ensemble import GradientBoostingClassifier
import xgboost as xgb
from xgboost import XGBClassifier
```

```
[75]: df = pd.read_csv('Customer Churn Prediction/WA_Fn-UseC_-Telco-Customer-Churn.
↳csv')
df
```

```
[75]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	
...	...	...	...	...	...	...	
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	
7040	4801-JZAZL	Female	0	Yes	Yes	11	
7041	8361-LTMKD	Male	1	Yes	No	4	

7042 3186-AJIEK Male 0 No No 66

	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\
0	No	No phone service	DSL	No	...	
1	Yes	No	DSL	Yes	...	
2	Yes	No	DSL	Yes	...	
3	No	No phone service	DSL	Yes	...	
4	Yes	No	Fiber optic	No	...	
...	...	...	...	...	...	
7038	Yes	Yes	DSL	Yes	...	
7039	Yes	Yes	Fiber optic	No	...	
7040	No	No phone service	DSL	Yes	...	
7041	Yes	Yes	Fiber optic	No	...	
7042	Yes	No	Fiber optic	Yes	...	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	
1	Yes	No	No	No	One year	
2	No	No	No	No	Month-to-month	
3	Yes	Yes	No	No	One year	
4	No	No	No	No	Month-to-month	
...	...	...	...	...	...	
7038	Yes	Yes	Yes	Yes	One year	
7039	Yes	No	Yes	Yes	One year	
7040	No	No	No	No	Month-to-month	
7041	No	No	No	No	Month-to-month	
7042	Yes	Yes	Yes	Yes	Two year	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
0	Yes	Electronic check	29.85	29.85	
1	No	Mailed check	56.95	1889.5	
2	Yes	Mailed check	53.85	108.15	
3	No	Bank transfer (automatic)	42.30	1840.75	
4	Yes	Electronic check	70.70	151.65	
...	...	...	...	...	
7038	Yes	Mailed check	84.80	1990.5	
7039	Yes	Credit card (automatic)	103.20	7362.9	
7040	Yes	Electronic check	29.60	346.45	
7041	Yes	Mailed check	74.40	306.6	
7042	Yes	Bank transfer (automatic)	105.65	6844.5	

	Churn
0	No
1	No
2	Yes
3	No
4	Yes

```
...
7038    No
7039    No
7040    No
7041    Yes
7042    No
```

```
[7043 rows x 21 columns]
```

```
[76]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService         7043 non-null   object
9   OnlineSecurity          7043 non-null   object
10  OnlineBackup            7043 non-null   object
11  DeviceProtection        7043 non-null   object
12  TechSupport             7043 non-null   object
13  StreamingTV             7043 non-null   object
14  StreamingMovies         7043 non-null   object
15  Contract                7043 non-null   object
16  PaperlessBilling        7043 non-null   object
17  PaymentMethod           7043 non-null   object
18  MonthlyCharges          7043 non-null   float64
19  TotalCharges            7043 non-null   object
20  Churn                   7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[77]: df.columns
```

```
[77]: Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
        'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
        'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
        'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
        'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
```

```
dtype='object')
```

## 2 Data Cleaning

```
[78]: # Turn total charges into a number value and check for missing values
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors = 'coerce')
df.isnull().sum()
```

```
[78]: customerID      0
gender            0
SeniorCitizen    0
Partner          0
Dependents       0
tenure           0
PhoneService     0
MultipleLines    0
InternetService  0
OnlineSecurity   0
OnlineBackup     0
DeviceProtection 0
TechSupport      0
StreamingTV      0
StreamingMovies  0
Contract         0
PaperlessBilling 0
PaymentMethod    0
MonthlyCharges   0
TotalCharges     11
Churn            0
dtype: int64
```

```
[79]: # Replace NaN with the column mean
df['TotalCharges'] = df['TotalCharges'].fillna(df['TotalCharges'].median())
df.isnull().sum()
```

```
[79]: customerID      0
gender            0
SeniorCitizen    0
Partner          0
Dependents       0
tenure           0
PhoneService     0
MultipleLines    0
InternetService  0
OnlineSecurity   0
OnlineBackup     0
```

```

DeviceProtection    0
TechSupport         0
StreamingTV         0
StreamingMovies     0
Contract            0
PaperlessBilling    0
PaymentMethod       0
MonthlyCharges      0
TotalCharges        0
Churn               0
dtype: int64

```

```
[80]: df = df.drop('customerID', axis = 1)
df
```

```
[80]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	
...	...	...	...	...	...	...	
7038	Male	0	Yes	Yes	24	Yes	
7039	Female	0	Yes	Yes	72	Yes	
7040	Female	0	Yes	Yes	11	No	
7041	Male	1	Yes	No	4	Yes	
7042	Male	0	No	No	66	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	No phone service	DSL	No	Yes	
1	No	DSL	Yes	No	
2	No	DSL	Yes	Yes	
3	No phone service	DSL	Yes	No	
4	No	Fiber optic	No	No	
...	...	...	...	...	
7038	Yes	DSL	Yes	No	
7039	Yes	Fiber optic	No	Yes	
7040	No phone service	DSL	Yes	No	
7041	Yes	Fiber optic	No	No	
7042	No	Fiber optic	Yes	No	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month-to-month	
1	Yes	No	No	No	One year	
2	No	No	No	No	Month-to-month	
3	Yes	Yes	No	No	One year	
4	No	No	No	No	Month-to-month	

...	...	...	...	...	...	...
7038	Yes	Yes	Yes	Yes	One year	
7039	Yes	No	Yes	Yes	One year	
7040	No	No	No	No	Month-to-month	
7041	No	No	No	No	Month-to-month	
7042	Yes	Yes	Yes	Yes	Two year	

	PaperlessBilling	PaymentMethod	MonthlyCharges \
0	Yes	Electronic check	29.85
1	No	Mailed check	56.95
2	Yes	Mailed check	53.85
3	No	Bank transfer (automatic)	42.30
4	Yes	Electronic check	70.70

...	...	...	...
7038	Yes	Mailed check	84.80
7039	Yes	Credit card (automatic)	103.20
7040	Yes	Electronic check	29.60
7041	Yes	Mailed check	74.40
7042	Yes	Bank transfer (automatic)	105.65

	TotalCharges	Churn
0	29.85	No
1	1889.50	No
2	108.15	Yes
3	1840.75	No
4	151.65	Yes

...	...	...
7038	1990.50	No
7039	7362.90	No
7040	346.45	No
7041	306.60	Yes
7042	6844.50	No

[7043 rows x 20 columns]

```
[81]: df.isnull().sum()
```

```
[81]: gender          0
      SeniorCitizen  0
      Partner        0
      Dependents     0
      tenure         0
      PhoneService   0
      MultipleLines   0
      InternetService 0
      OnlineSecurity  0
      OnlineBackup    0
```

```

DeviceProtection    0
TechSupport         0
StreamingTV         0
StreamingMovies     0
Contract            0
PaperlessBilling    0
PaymentMethod       0
MonthlyCharges      0
TotalCharges        0
Churn               0
dtype: int64

```

```
[82]: df.describe()
```

```

[82]:      SeniorCitizen    tenure  MonthlyCharges  TotalCharges
count      7043.000000  7043.000000    7043.000000    7043.000000
mean         0.162147    32.371149     64.761692    2281.916928
std          0.368612    24.559481     30.090047    2265.270398
min          0.000000     0.000000     18.250000     18.800000
25%          0.000000     9.000000     35.500000    402.225000
50%          0.000000    29.000000     70.350000   1397.475000
75%          0.000000    55.000000     89.850000   3786.600000
max          1.000000    72.000000    118.750000   8684.800000

```

### 3 Data Visualizations

```

[83]: # Frequency of Churn/No Churn
churn_stats = df['Churn'].value_counts(normalize=True) * 100
churn_stats

```

```

[83]: Churn
No      73.463013
Yes     26.536987
Name: proportion, dtype: float64

```


```

[84]: # Data visualization to churn/no churn

labels = 'Churn', 'No Churn'
sizes = [26.536987, 73.463013]

fig, ax = plt.subplots()
ax.pie(sizes, labels=labels,
      wedgeprops=dict(edgecolor='white'), autopct='%1.1f%%')
plt.show()

```



```
[85]: # Calculate the average of monthly charges for no churn
no_churn = df[df['Churn'] == 'No']
average_monthly_no_churn = no_churn['MonthlyCharges'].mean()
print(f"Average Monthly Charges for No Churn: {average_monthly_no_churn}")

# Calculate the median of monthly charges for no churn
median_monthly_no_churn = no_churn['MonthlyCharges'].median()
print(f"Median Monthly Charges for No Churn: {median_monthly_no_churn}")

# Calculate the average of monthly charges for churn
yes_churn = df[df['Churn'] == 'Yes']
average_monthly_yes_churn = yes_churn['MonthlyCharges'].mean()
print(f"Average Monthly Charges for Churn: {average_monthly_yes_churn}")

# Calculate the median of monthly charges for churn
median_monthly_yes_churn = yes_churn['MonthlyCharges'].median()
print(f"Median Monthly Charges for Churn: {median_monthly_yes_churn}")
```

```
Average Monthly Charges for No Churn: 61.26512369540008
Median Monthly Charges for No Churn: 64.42500000000001
Average Monthly Charges for Churn: 74.44133226324237
Median Monthly Charges for Churn: 79.65
```




```
[86]: # Data visualization of churn/no churn based on monthly charges

# Filter and plot the data for 'Churn' == 0 and 'Churn' == 1
sns.kdeplot(df.MonthlyCharges[df["Churn"] == 'No'], fill = True, label="No Churn")
sns.kdeplot(df.MonthlyCharges[df["Churn"] == 'Yes'], fill = True, label="Churn")

# Add labels and title
plt.title('Monthly Charges by Churn (KDE PLOT)')
plt.xlabel('Monthly Charges')
plt.ylabel('Density')

# Display the legend and the plot
plt.show()
```



```
[87]: # Data visualization of churn/no churn based on monthly charges (histogram)

# Filter and plot the data for 'Churn' == 0 and 'Churn' == 1
sns.histplot(df.MonthlyCharges[df["Churn"] == 'No'], bins = 30, alpha = 0.5, label="No Churn")
```


```

sns.histplot(df.MonthlyCharges[df["Churn"] == 'Yes'], bins = 30, alpha = 0.5,
             label="Churn")

# Add labels and title
plt.title('Monthly Charges by Churn (HISTOGRAM)')
plt.xlabel('Monthly Charges')
plt.ylabel('Density')

# Display the legend and the plot
plt.show()

```



```

[88]: # Calculate the average of total charges for no churn
average_total_no_churn = no_churn['TotalCharges'].mean()
print(f"Average Total Charges for No Churn : {average_total_no_churn}")

# Calculate the average of total charges for churn
average_total_yes_churn = yes_churn['TotalCharges'].mean()
print(f"Average Total Charges for Churn: {average_total_yes_churn}")

# Calculate the median of total charges for no churn

```

```

median_total_no_churn = no_churn['TotalCharges'].median()
print(f"Median Total Charges for No Churn: {median_total_no_churn}")

# Calculate the median of total charges for churn
median_total_yes_churn = yes_churn['TotalCharges'].median()
print(f"Median Total Charges for Churn: {median_total_yes_churn}")

```

Average Total Charges for No Churn : 2552.882494201778  
 Average Total Charges for Churn: 1531.7960941680042  
 Median Total Charges for No Churn: 1679.525  
 Median Total Charges for Churn: 703.55

```


[89]: # Data visualization of churn/no churn based on total charges (kde plot)

# Filter and plot the data for 'Churn' == 0 and 'Churn' == 1
sns.kdeplot(df.TotalCharges[df["Churn"] == 'No'], fill = True, label="No Churn")
sns.kdeplot(df.TotalCharges[df["Churn"] == 'Yes'], fill = True, label="Churn")

# Add labels and title
plt.title('Total Charges by Churn (KDE PLOT)')
plt.xlabel('Total Charges')
plt.ylabel('Density')

# Display the legend and the plot
plt.show()

```

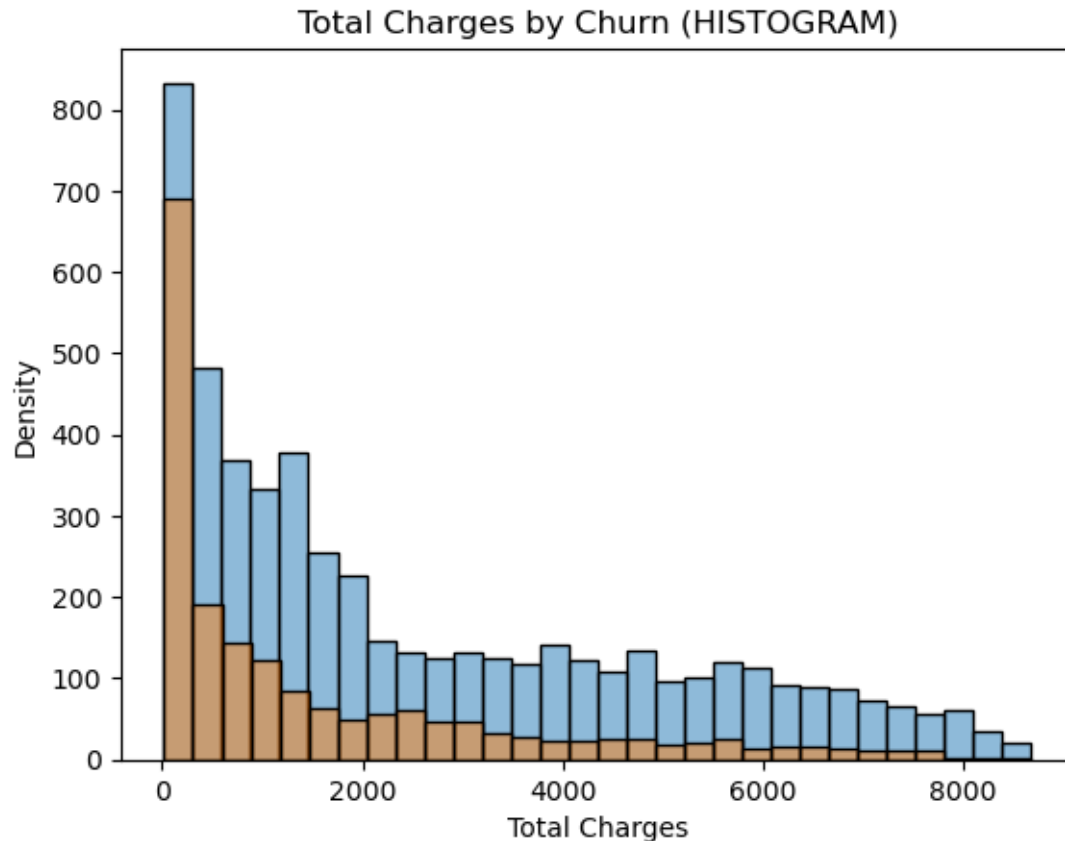


```
[90]: # Data visualization of churn/no churn based on total charges

# Filter and plot the data for 'Churn' == 0 and 'Churn' == 1
sns.histplot(df.TotalCharges[df["Churn"] == 'No'], bins = 30, alpha = 0.5,
             label="No Churn")
sns.histplot(df.TotalCharges[df["Churn"] == 'Yes'], bins = 30, alpha = 0.5,
             label="Churn")

# Add labels and title
plt.title('Total Charges by Churn (HISTOGRAM)')
plt.xlabel('Total Charges')
plt.ylabel('Density')

# Display the legend and the plot
plt.show()
```



```
[91]: # Calculate the average tenure for no churn
average_tenure_no_churn = no_churn['tenure'].mean()
print(f"Average Tenure for No Churn : {average_tenure_no_churn}")

# Calculate the median tenure for no churn
median_tenure_no_churn = no_churn['tenure'].median()
print(f"Median Tenure for No Churn: {median_tenure_no_churn}")

# Calculate the average tenure for churn
average_tenure_yes_churn = yes_churn['tenure'].mean()
print(f"Average Tenure for Churn: {average_tenure_yes_churn}")


# Calculate the median tenure for churn
median_tenure_yes_churn = yes_churn['tenure'].median()
print(f"Median Tenure for Churn: {median_tenure_yes_churn}")
```

```
Average Tenure for No Churn : 37.56996521066873
Median Tenure for No Churn: 38.0
Average Tenure for Churn: 17.979133226324237
Median Tenure for Churn: 10.0
```

```
[92]: # Data visualization for tenure by churn
sns.boxplot(data = df, x = 'Churn', y = 'tenure')

# Add labels and title
plt.title('Tenure by Churn')
plt.xlabel('Churn')
plt.ylabel('Tenure')

# Display the legend and the plot
plt.show()
```




```
[93]: # Data visualization for tenure by churn
sns.violinplot(data = df, x = 'Churn', y = 'tenure')

# Add labels and title
plt.title('Tenure by Churn')
plt.xlabel('Churn')
plt.ylabel('Tenure')

# Display the legend and the plot
```

```
plt.show()
```



```
[94]: # Percentage of Churn vs. No Churn

# Count occurrences of churn and non-churn
churn_counts = df['Churn'].value_counts()

# Calculate percentage
churn_percentage = (churn_counts / len(df)) * 100
churn_percentage
```

```
[94]: Churn
No      73.463013
Yes     26.536987
Name: count, dtype: float64
```

```
[95]: # Data visualization for Contract Type by Churn

#Frequency Table for contract type
```

```

contracttype_churn_counts = df.groupby(['Churn', 'Contract']).size().
    ↪unstack(fill_value=0)
print(contracttype_churn_counts)
print("\n")
# Normalized frequency table for the contract type
contract_table_percent = contracttype_churn_counts.
    ↪div(contracttype_churn_counts.sum(axis=1), axis=0) * 100
print(contract_table_percent)

# Create bar chart

# Plot the bar chart
contracttype_churn_counts.plot(kind='bar', figsize=(8, 6))
plt.title('Contract Type by Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.legend(['Month-to-Month', 'One Year', 'Two Year'], loc='upper right')

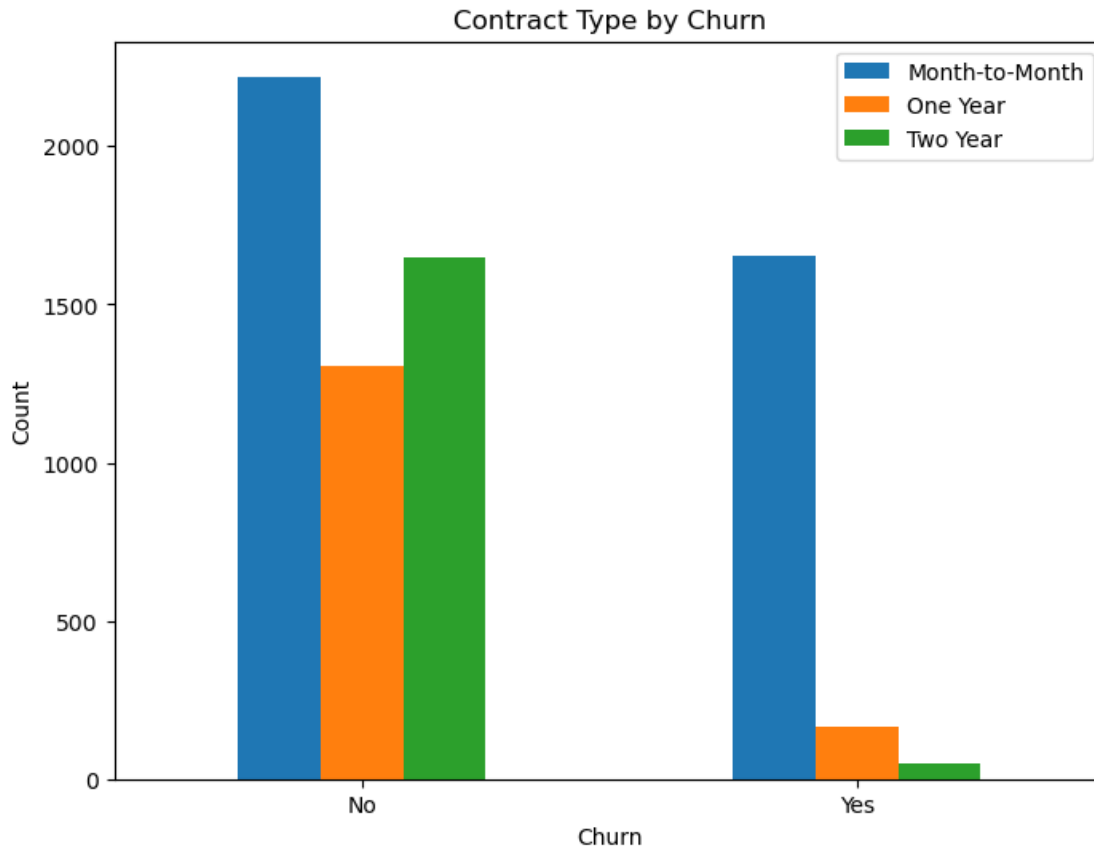
# Show the chart
plt.show()

```

Contract	Month-to-month	One year	Two year
Churn			
No	2220	1307	1647
Yes	1655	166	48

Contract	Month-to-month	One year	Two year
Churn			
No	42.906842	25.260920	31.832238
Yes	88.550027	8.881755	2.568218





```
[96]: # Frequency for each type of payment method
payment_stats = df['PaymentMethod'].value_counts(normalize=True) * 100
payment_stats
```


```
[96]: PaymentMethod
Electronic check      33.579441
Mailed check         22.887974
Bank transfer (automatic)  21.922476
Credit card (automatic)  21.610109
Name: proportion, dtype: float64
```

```
[97]: # Data visualization to show disitrbution of all four payment methods

labels = 'Electronic check', 'Mailed check', 'Bank transfer (automatic)', 'Credit card'
sizes = [33.579441, 22.887974, 21.922476, 21.610109]

fig, ax = plt.subplots()
ax.pie(sizes, labels=labels,
```

```
wedgeprops=dict(edgecolor='white'), autopct='%1.1f%%')
plt.show()
```



```
[98]: payment_dist = df.groupby('Churn')['PaymentMethod'].value_counts().unstack()
print(payment_dist)
print('\n')

# Percentage for Payment Method in Customer Churn overall
payment_churn_percentage = df.groupby('Churn')['PaymentMethod'].
    value_counts(normalize=True).mul(100).unstack()
print(payment_churn_percentage)
```

PaymentMethod	Bank transfer (automatic)	Credit card (automatic)	\
Churn			
No	1286	1290	
Yes	258	232	

PaymentMethod	Electronic check	Mailed check
Churn		
No	1294	1304
Yes	1071	308

PaymentMethod	Bank transfer (automatic)	Credit card (automatic)	\
Churn			
No	24.855044	24.932354	
Yes	13.804173	12.413055	

PaymentMethod	Electronic check	Mailed check
Churn		
No	25.009664	25.202938
Yes	57.303371	16.479401

```
[99]: # Data visualization for payment method by Churn

#Frequency Table for contract type
payment_churn_counts = df.groupby(['Churn', 'PaymentMethod']).size().
    ↳unstack(fill_value=0)
print(payment_churn_counts)
print("\n")

# Normalized frequency table for the contract type
payment_table_percent = payment_churn_counts.div(payment_churn_counts.
    ↳sum(axis=1), axis=0) * 100
print(payment_table_percent)

# Create bar chart

# Plot the bar chart
payment_churn_counts.plot(kind='bar', figsize=(8, 6))
plt.title('Payment Method by Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.legend(['Bank transfer (automatic)', 'Credit card', 'Electronic check',
    ↳'Mailed check'], loc='upper right')

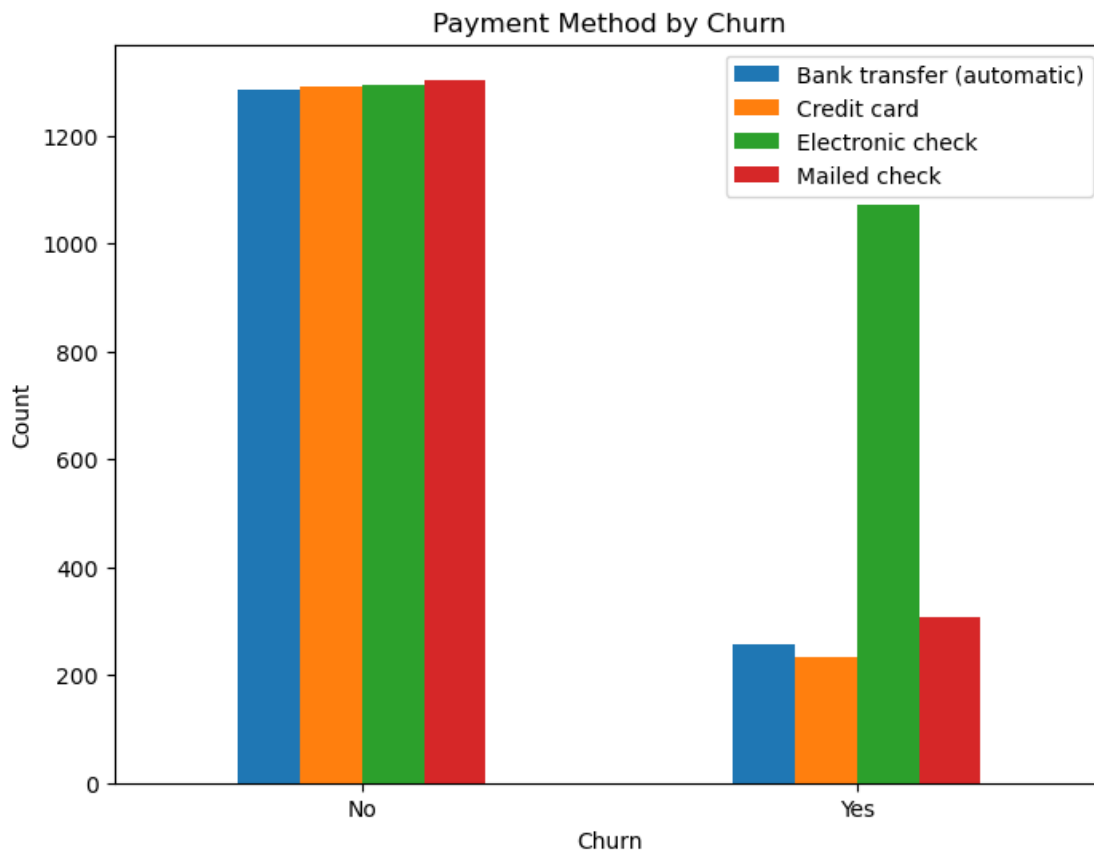
# Show the chart
plt.show()
```

PaymentMethod	Bank transfer (automatic)	Credit card (automatic)	\
Churn			
No	1286	1290	
Yes	258	232	

PaymentMethod	Electronic check	Mailed check
Churn		
No	1294	1304
Yes	1071	308

PaymentMethod	Bank transfer (automatic)	Credit card (automatic)	\
Churn			
No	24.855044	24.932354	
Yes	13.804173	12.413055	

PaymentMethod	Electronic check	Mailed check
Churn		
No	25.009664	25.202938
Yes	57.303371	16.479401



```
[100]: df = df.rename(columns={
    'PhoneService': 'Phone Service',
    'MultipleLines': 'Multiple Lines',
    'InternetService': 'Internet Service',
    'OnlineSecurity': 'Online Security',
    'OnlineBackup': 'Online Backup',
    'DeviceProtection': 'Device Protection',
    'TechSupport': 'Tech Support',
    'StreamingTV': 'Streaming TV',
})
```

```

        'StreamingMovies': 'Streaming Movies',
    })

service = ['Phone Service', 'Multiple Lines', 'Internet Service', 'Online_
↳Security',
           'Online Backup', 'Device Protection', 'Tech Support', 'Streaming_
↳TV', 'Streaming Movies']
def generate_service_frequency_by_churn(df, service):
    for col in service:
        print(f"\nFrequency Table for '{col}' (Grouped by Churn):")
        print(df.groupby('Churn')[col].value_counts()) # Raw counts
        print("\nPercentage Distribution by Churn:")
        print(df.groupby('Churn')[col].value_counts(normalize=True).mul(100).
↳round(2)) # Percentage
        print("-" * 60)

# Call function on the service columns
generate_service_frequency_by_churn(df, service)

```

Frequency Table for 'Phone Service' (Grouped by Churn):

Churn	Phone Service	
No	Yes	4662
	No	512
Yes	Yes	1699
	No	170

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Phone Service	
No	Yes	90.1
	No	9.9
Yes	Yes	90.9
	No	9.1

Name: proportion, dtype: float64

Frequency Table for 'Multiple Lines' (Grouped by Churn):

Churn	Multiple Lines	
No	No	2541
	Yes	2121
	No phone service	512
Yes	Yes	850
	No	849
	No phone service	170

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn Multiple Lines

No	No	49.11
	Yes	40.99
	No phone service	9.90
Yes	Yes	45.48
	No	45.43
	No phone service	9.10

Name: proportion, dtype: float64

-----

Frequency Table for 'Internet Service' (Grouped by Churn):

Churn Internet Service

No	DSL	1962
	Fiber optic	1799
	No	1413
Yes	Fiber optic	1297
	DSL	459
	No	113

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn Internet Service

No	DSL	37.92
	Fiber optic	34.77
	No	27.31
Yes	Fiber optic	69.40
	DSL	24.56
	No	6.05

Name: proportion, dtype: float64

-----

Frequency Table for 'Online Security' (Grouped by Churn):

Churn Online Security

No	No	2037
	Yes	1724
	No internet service	1413
Yes	No	1461
	Yes	295
	No internet service	113

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn Online Security

No	No	39.37
	Yes	33.32
	No internet service	27.31
Yes	No	78.17

```

        Yes                15.78
        No internet service  6.05
Name: proportion, dtype: float64
-----
```

Frequency Table for 'Online Backup' (Grouped by Churn):

```
Churn  Online Backup
No      Yes                1906
        No                1855
        No internet service 1413
Yes     No                1233
        Yes                523
        No internet service 113
Name: count, dtype: int64
```

Percentage Distribution by Churn:

```
Churn  Online Backup
No      Yes                36.84
        No                35.85
        No internet service 27.31
Yes     No                65.97
        Yes                27.98
        No internet service  6.05
Name: proportion, dtype: float64
-----
```

Frequency Table for 'Device Protection' (Grouped by Churn):

```
Churn  Device Protection
No      No                1884
        Yes                1877
        No internet service 1413
Yes     No                1211
        Yes                545
        No internet service 113
Name: count, dtype: int64
```

Percentage Distribution by Churn:

```
Churn  Device Protection
No      No                36.41
        Yes                36.28
        No internet service 27.31
Yes     No                64.79
        Yes                29.16
        No internet service  6.05
Name: proportion, dtype: float64
-----
```

Frequency Table for 'Tech Support' (Grouped by Churn):

Churn	Tech Support	
No	No	2027
	Yes	1734
	No internet service	1413
Yes	No	1446
	Yes	310
	No internet service	113

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Tech Support	
No	No	39.18
	Yes	33.51
	No internet service	27.31
Yes	No	77.37
	Yes	16.59
	No internet service	6.05

Name: proportion, dtype: float64

Frequency Table for 'Streaming TV' (Grouped by Churn):

Churn	Streaming TV	
No	Yes	1893
	No	1868
	No internet service	1413
Yes	No	942
	Yes	814
	No internet service	113

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Streaming TV	
No	Yes	36.59
	No	36.10
	No internet service	27.31
Yes	No	50.40
	Yes	43.55
	No internet service	6.05

Name: proportion, dtype: float64

Frequency Table for 'Streaming Movies' (Grouped by Churn):

Churn	Streaming Movies	
No	Yes	1914
	No	1847
	No internet service	1413
Yes	No	938
	Yes	818



```
      No internet service      113
Name: count, dtype: int64
```

Percentage Distribution by Churn:

```
Churn  Streaming Movies
No      Yes              36.99
        No              35.70
        No internet service  27.31
Yes      No              50.19
        Yes              43.77
        No internet service   6.05
```

```
Name: proportion, dtype: float64
```


-----

```
[101]: # Data visualization for different telecommunication services by Churn

# Set up the figure and axes
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(15, 12)) # 3x3 grid
axes = axes.flatten() # Flatten to loop easily

# Generate bar charts for each service column
for i, col in enumerate(service):
    sns.countplot(data=df, x='Churn', hue=col, ax=axes[i])
    axes[i].set_title(f"{col} by Churn")
    axes[i].set_xlabel("Churn")
    axes[i].set_ylabel("Count")
    axes[i].tick_params(axis='x', rotation=45) # Rotate x-axis labels for
    ↪readability

# Adjust layout to avoid overlap
plt.tight_layout()
plt.show()
```



```
[102]: # Create frequency table for each demographic
```

```
# Rename columns
```

```
df = df.rename(columns = {
    'gender' : 'Gender',
    'SeniorCitizen' : 'Senior Citizen'
})
```

```
# Rename values in the 'Senior Citizen' column
```

```
df['Senior Citizen'] = df['Senior Citizen'].replace({0: 'No', 1: 'Yes'})
```

```
# Data visualization for demographics
```

```
demographics = ['Gender', 'Senior Citizen', 'Partner', 'Dependents']
```

```
# Frequency table function
```

```
def generate_demographic_frequency_by_churn(df, demographics):
    for col in demographics:
        print(f"\nFrequency Table for '{col}' (Grouped by Churn):")
        print(df.groupby('Churn')[col].value_counts()) # Raw counts
```

```

        print("\nPercentage Distribution by Churn:")
        print(df.groupby('Churn')[col].value_counts(normalize=True).mul(100).
↳round(2)) # Percentage
        print("-" * 60)

# Call function on the demographics columns
generate_service_frequency_by_churn(df, demographics)

```

Frequency Table for 'Gender' (Grouped by Churn):

Churn	Gender	
No	Male	2625
	Female	2549
Yes	Female	939
	Male	930

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Gender	
No	Male	50.73
	Female	49.27
Yes	Female	50.24
	Male	49.76

Name: proportion, dtype: float64

-----

Frequency Table for 'Senior Citizen' (Grouped by Churn):

Churn	Senior Citizen	
No	No	4508
	Yes	666
Yes	No	1393
	Yes	476

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Senior Citizen	
No	No	87.13
	Yes	12.87
Yes	No	74.53
	Yes	25.47

Name: proportion, dtype: float64

-----

Frequency Table for 'Partner' (Grouped by Churn):

Churn	Partner	
No	Yes	2733
	No	2441

Yes	No	1200
	Yes	669

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Partner	
No	Yes	52.82
	No	47.18
Yes	No	64.21
	Yes	35.79

Name: proportion, dtype: float64

-----

Frequency Table for 'Dependents' (Grouped by Churn):

Churn	Dependents	
No	No	3390
	Yes	1784
Yes	No	1543
	Yes	326

Name: count, dtype: int64

Percentage Distribution by Churn:

Churn	Dependents	
No	No	65.52
	Yes	34.48
Yes	No	82.56
	Yes	17.44

Name: proportion, dtype: float64

-----

```
[103]: # Set up the figure (for a 2x2 grid layout)
fig, axes = plt.subplots(2, 2, figsize=(10, 8)) # 2x2 grid for 4 demographic
        ↪ features
axes = axes.flatten()

# Loop through demographic columns to create pie charts
for i, col in enumerate(demographics):
    ax = axes[i]

    # Get the value counts for the demographic column
    demographic_counts = df[col].value_counts()


    # Create pie chart for the demographic distribution
    ax.pie(
        demographic_counts, labels=demographic_counts.index, autopct='%1.1f%%',
        ↪ startangle=90,
        colors=['skyblue', 'salmon'], wedgeprops={'edgecolor': 'white'})
```

```

    )
    ax.set_title(f'{col} Distribution')

# Adjust layout and display the pie charts
plt.tight_layout()
plt.show()

```



```

[104]: # Set up the figure and axes
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(15, 12)) # 2x2 grid
axes = axes.flatten() # Flatten to loop easily

# Generate bar charts for each service column
for i, col in enumerate(demographics):
    sns.countplot(data=df, x='Churn', hue=col, ax=axes[i])
    axes[i].set_title(f'{col} by Churn')


```

```

axes[i].set_xlabel("Churn")
axes[i].set_ylabel("Count")
axes[i].tick_params(axis='x', rotation=0) # Rotate x-axis labels for
↳ readability

# Adjust layout to avoid overlap
plt.tight_layout()
plt.show()

```



[105]: # Data visualization for paperless billing by Churn

```

#Frequency Table for contract type
billing_churn_counts = df.groupby(['Churn', 'PaperlessBilling']).size().
↳unstack(fill_value=0)
print(billing_churn_counts)
print("\n")

# Normalized frequency table for the contract type

```

```

billing_table_percent = billing_churn_counts.div(billing_churn_counts.
    ↪sum(axis=1), axis=0) * 100
print(billing_table_percent)

# Create bar chart

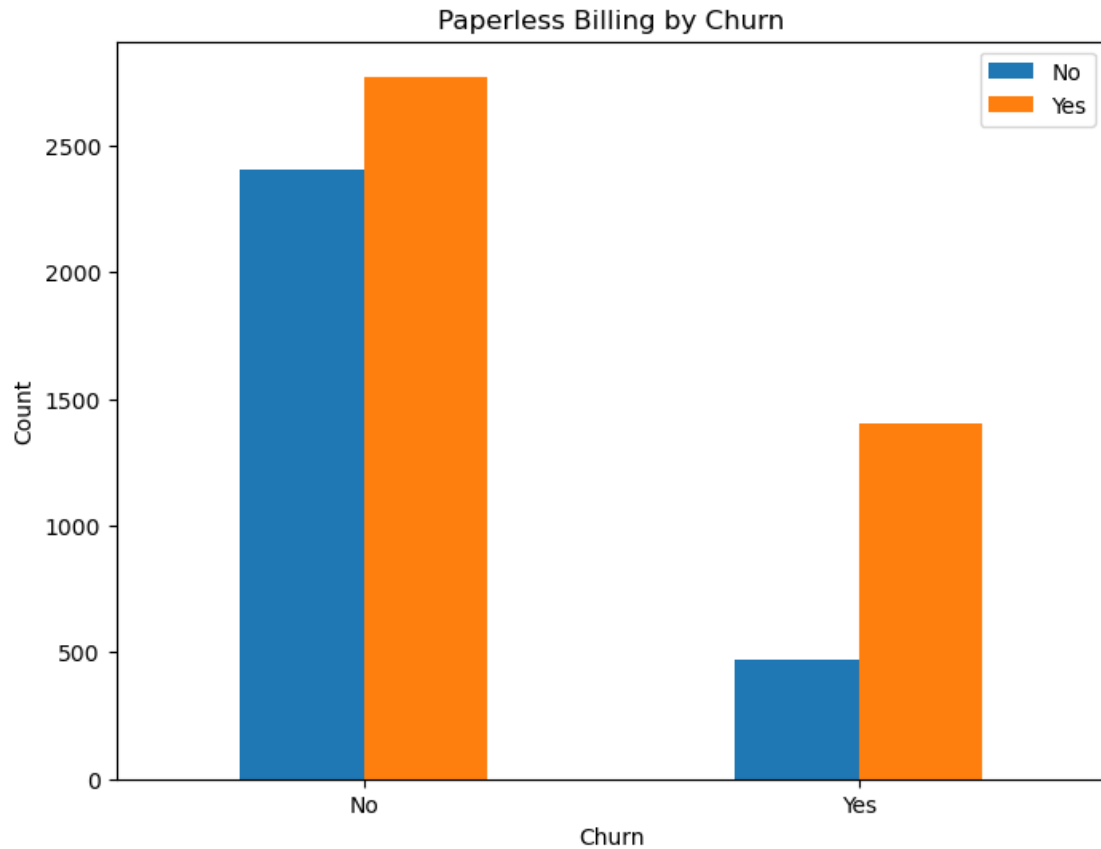
# Plot the bar chart
billing_churn_counts.plot(kind='bar', figsize=(8, 6))
plt.title('Paperless Billing by Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.legend(['No', 'Yes'], loc='upper right')

# Show the chart
plt.show()

```

PaperlessBilling	No	Yes
Churn		
No	2403	2771
Yes	469	1400

PaperlessBilling	No	Yes
Churn		
No	46.443757	53.556243
Yes	25.093633	74.906367



## 4 Data Processing

```
[106]: # Turn categorical variables into numbers
df_copy = df.copy()
columns = df_copy.columns
label_encoder = LabelEncoder()
for col in columns:
    df_copy[col] = label_encoder.fit_transform(df_copy[col])
df_copy
```

```
[106]:
```

	Gender	Senior Citizen	Partner	Dependents	tenure	Phone Service	\
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	
...	...	...	...	...	...	...	
7038	1	0	1	1	24	1	
7039	0	0	1	1	72	1	



7040	0	0	1	1	11	0
7041	1	1	1	0	4	1
7042	1	0	0	0	66	1

	Multiple Lines	Internet Service	Online Security	Online Backup	\
0	1	0	0	2	
1	0	0	2	0	
2	0	0	2	2	
3	1	0	2	0	
4	0	1	0	0	
...	...	...	...	...	
7038	2	0	2	0	
7039	2	1	0	2	
7040	1	0	2	0	
7041	2	1	0	0	
7042	0	1	2	0	

	Device Protection	Tech Support	Streaming TV	Streaming Movies	\
0	0	0	0	0	
1	2	0	0	0	
2	0	0	0	0	
3	2	2	0	0	
4	0	0	0	0	
...	...	...	...	...	
7038	2	2	2	2	
7039	2	0	2	2	
7040	0	0	0	0	
7041	0	0	0	0	
7042	2	2	2	2	

	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\
0	0	1	2	142	74	
1	1	0	3	498	3625	
2	0	1	3	436	536	
3	1	0	0	266	3571	
4	0	1	2	729	674	
...	...	...	...	...	...	
7038	1	1	3	991	3701	
7039	1	1	1	1340	6305	
7040	0	1	2	137	1265	
7041	0	1	3	795	1157	
7042	2	1	0	1388	6151	

Churn	
0	0
1	0
2	1

```

3          0
4          1
...      ...
7038       0
7039       0
7040       0
7041       1
7042       0

```

```
[7043 rows x 20 columns]
```


```
[107]: # Find the correlation between Churn and other variables
correlation_matrix = df_copy.corr()
churn_correlation = correlation_matrix['Churn'].sort_values(ascending=False)
print(churn_correlation)
```

```

Churn          1.000000
PaperlessBilling  0.191825
MonthlyCharges  0.183523
Senior Citizen  0.150889
PaymentMethod   0.107062
Multiple Lines  0.038037
Phone Service   0.011942
Gender          -0.008612
Streaming TV    -0.036581
Streaming Movies -0.038492
Internet Service -0.047291
Partner         -0.150448
Dependents      -0.164221
Device Protection -0.178134
Online Backup    -0.195525
TotalCharges    -0.230754
Tech Support    -0.282492
Online Security -0.289309
tenure          -0.352229
Contract        -0.396713
Name: Churn, dtype: float64

```

```
[108]: # Visualize with a heat map
plt.figure(figsize=(15, 12)) # Adjust figure size
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5,
            vmin = -1, vmax = 1)
plt.title('Feature Correlation Heatmap')
plt.show()
```



```
[109]: # Define features and target
X = df_copy.drop(columns=['Churn'])
y = df_copy['Churn']
X = pd.get_dummies(X, drop_first=True)

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

## 5 Machine Learning Models and Predictions

```
[110]: # Train with Random Forest
random_forest = RandomForestClassifier(n_estimators=100, random_state=42)
random_forest.fit(X_train, y_train)
y_pred_RF = random_forest.predict(X_test)
```

```
# Evaluate performance
accuracy = accuracy_score(y_test, y_pred_RF)
print(f"Accuracy: {accuracy * 100}%")
```

Accuracy: 79.63094393186657%

```
[111]: # Confusion matrix for Random Forest
```

```
confusion_matrix_RF = metrics.confusion_matrix(y_test, y_pred)
confusion_matrix_RF
```

```
[111]: array([[946,  90],
              [197, 176]])
```

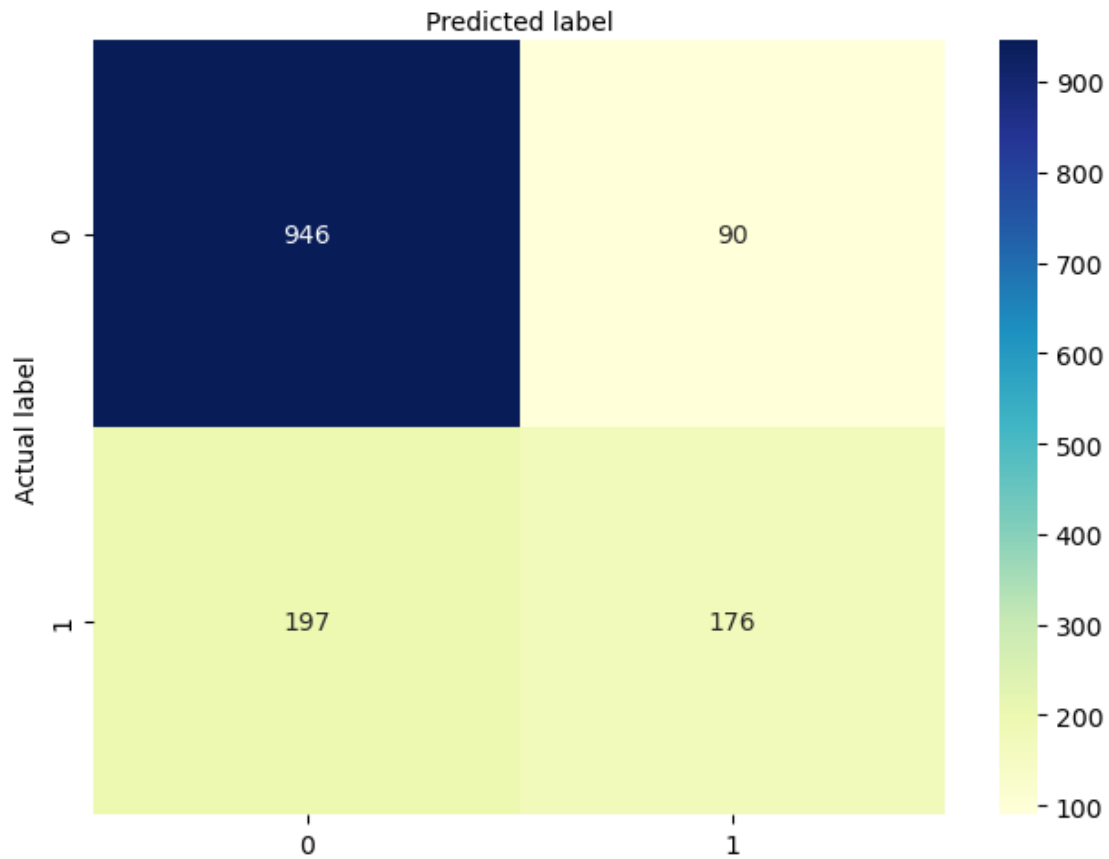
```
[112]: # Create confusion matrix heat map for Logistic Regression
```

```
class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)

# create heatmap
sns.heatmap(pd.DataFrame(confusion_matrix_RF), annot=True, cmap="YlGnBu",
            ↪,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion Matrix for Random Forest Classifier', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
[112]: Text(0.5, 427.95555555555555, 'Predicted label')
```

Confusion Matrix for Random Forest Classifier




```
[113]: target_names = ['No Churn', 'Churn']
print(classification_report(y_test, y_pred_RF, target_names=target_names))
```

	precision	recall	f1-score	support
No Churn	0.83	0.91	0.87	1036
Churn	0.66	0.47	0.55	373
accuracy			0.80	1409
macro avg	0.74	0.69	0.71	1409
weighted avg	0.78	0.80	0.78	1409

```
[114]: # Predict probabilities
y_probs = random_forest.predict_proba(X_test)[:, 1]
fpr_rf, tpr_rf, thresholds = roc_curve(y_test, y_probs)
plt.plot([0, 1], [0, 1], 'k--' )
```

```
plt.plot(fpr_rf, tpr_rf, label='Random Forest',color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve',fontsize=16)
plt.show();
```



```
[115]: # Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train logistic regression model
logisticRegr = LogisticRegression(max_iter=500)
logisticRegr.fit(X_train_scaled, y_train)

# Model evaluation
accuracy = logisticRegr.score(X_test_scaled, y_test)
y_pred_log = logisticRegr.predict(X_test_scaled)
print(f"Accuracy: {accuracy * 100}%")
```

Accuracy: 81.33427963094393%

```
[116]: # Create a regression table

# Fit logistic regression model
logit_model = sm.Logit(y, X)
result = logit_model.fit()

# Display the regression table
print(result.summary())
```

Optimization terminated successfully.

Current function value: 0.412298

Iterations 7

#### Logit Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          7043
Model:                  Logit    Df Residuals:              7024
Method:                 MLE      Df Model:                 18
Date:                   Thu, 03 Apr 2025    Pseudo R-squ.:         0.2874
Time:                   16:39:00    Log-Likelihood:        -2903.8
converged:              True      LL-Null:              -4075.1
Covariance Type:        nonrobust    LLR p-value:           0.000
=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
Gender          -0.0481      0.064      -0.755      0.450      -0.173
0.077
Senior Citizen    0.2351      0.085       2.778      0.005       0.069
0.401
Partner          0.0273      0.078       0.350      0.727      -0.126
0.180
Dependents       -0.1769      0.090      -1.974      0.048      -0.353
-0.001
tenure           0.0114      0.006       1.912      0.056      -0.000
0.023
Phone Service    -1.0173      0.118      -8.622      0.000      -1.249
-0.786
Multiple Lines    0.0881      0.041       2.142      0.032       0.008
0.169
Internet Service  0.0946      0.064       1.479      0.139      -0.031
0.220
Online Security  -0.2578      0.041      -6.224      0.000      -0.339
-0.177
Online Backup     -0.1416      0.038      -3.700      0.000      -0.217
```

-0.067					
Device Protection	-0.0739	0.040	-1.864	0.062	-0.152
0.004					
Tech Support	-0.2465	0.042	-5.848	0.000	-0.329
-0.164					
Streaming TV	-0.0054	0.042	-0.129	0.897	-0.088
0.077					
Streaming Movies	-0.0005	0.042	-0.012	0.991	-0.082
0.081					
Contract	-0.8389	0.079	-10.585	0.000	-0.994
-0.684					
PaperlessBilling	0.3209	0.071	4.504	0.000	0.181
0.460					
PaymentMethod	0.0110	0.031	0.360	0.719	-0.049
0.071					
MonthlyCharges	0.0027	0.000	16.364	0.000	0.002
0.003					
TotalCharges	-0.0006	7.79e-05	-8.299	0.000	-0.001
-0.000					

```
=====
=====
```

```
[117]: # Evaluate classification model for Logistical Regression Model
```

```
confusion_matrix = metrics.confusion_matrix(y_test, y_pred_log)
confusion_matrix
```

```
[117]: array([[938,  98],
              [165, 208]])
```

```
[118]: # Create confusion matrix heat map for Logistic Regression
```


```
class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)

# create heatmap
sns.heatmap(pd.DataFrame(confusion_matrix), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix for Logistical Regression', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
[118]: Text(0.5, 427.9555555555555, 'Predicted label')
```



Confusion matrix for Logistical Regression



```
[119]: target_names = ['No Churn', 'Churn']
print(classification_report(y_test, y_pred_log, target_names=target_names))
```

	precision	recall	f1-score	support
No Churn	0.85	0.91	0.88	1036
Churn	0.68	0.56	0.61	373
accuracy			0.81	1409
macro avg	0.77	0.73	0.74	1409
weighted avg	0.81	0.81	0.81	1409

```
[ ]:
```